

1. Two Layer NN

a. Forward Pass:

i. Forward pass:

```
def relu(a):
    result = a * (a > 0)
    return result

#First-layer output
l1_output = np.dot(X, W1.T) + b1
l1_relu = relu(l1_output)

#Second-layer output
l2_output = np.dot(l1_relu, W2.T) + b2
#Using softmax as the output score
scores = l2_output

# ===== #
# END YOUR CODE HERE
# ===== #

# If the targets are not given then jump out, we're done
if y is None:
    return scores
```

ii.

iii. Loss Calculation:

```
# scores is num_examples by num_classes
def l2_norm(a):
    result = np.sqrt(np.sum(a**2))
    return result
def softmax(y_hat):
    exps = np.exp(y_hat - np.max(y_hat, axis=1, keepdims=True))
    result = exps / np.sum(exps)
    return result
def softmax_loss(sc, y_actual):
    sc = softmax(sc)
    yhat = sc[np.arange(sc.shape[0]), y_actual]
    result = -np.sum(np.log(yhat)) / sc.shape[0]
    return result

# Calculate the l2 norm for each weight
W1_l2norm = l2_norm(W1)
W2_l2norm = l2_norm(W2)
reg_loss = reg * (W1_l2norm + W2_l2norm)

loss = softmax_loss(scores, y) + reg_loss
# ===== #
# END YOUR CODE HERE
# ===== #
```

iv.

b. Gradient:

```
def relu_backward(dout, cache):
    x = cache
    dx = dout * (x > 0)
    return dx
# Calculate dW2
dW2 = (relu(np.dot(X, W1.T) + b1).T @ np.ones(scores.shape)).T
# Calculate db2 (derivative of softmax)
# Back pass value of b2 is only one
db2 = 1 #np.ones(b2.shape)
# Calculate dW1
dW1 = X.T @ relu_backward(np.ones(scores.shape) @ W2, np.dot(X, W1.T) + b1)
dW1 = dW1.T
# Calculate db1
db1 = np.mean(relu_backward(np.ones(scores.shape) @ W2, np.dot(X, W1.T) + b1).T, axis=1)

#print("Finish Calculation")

grads['W1'] = dW1
grads['b1'] = db1
grads['W2'] = dW2
grads['b2'] = db2
```

i.

c. Minibatch and learning rate:

```

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    # ===== #
    # YOUR CODE HERE:
    # Create a minibatch by sampling batch_size samples randomly.
    # ===== #
    random_data_index = np.random.randint(num_train, size=batch_size)
    X_batch = X[random_data_index, :]
    y_batch = y[random_data_index]
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # Compute loss and gradients using the current minibatch
    loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
    loss_history.append(loss)

    # ===== #
    # YOUR CODE HERE:
    # Perform a gradient descent step using the minibatch to update
    # all parameters (i.e., W1, W2, b1, and b2).
    # ===== #

    self.params['W1'] -= learning_rate * grads['W1']
    self.params['b1'] -= learning_rate * grads['b1']
    self.params['W2'] -= learning_rate * grads['W2']
    self.params['b2'] -= learning_rate * grads['b2']

```

i.

d. Prediction:

i.

```

"""
y_pred = None

# ===== #
# YOUR CODE HERE:
# Predict the class given the input data.
# ===== #

def relu(a):
    result = np.maximum(np.zeros(a.shape), a)
    return result
l1_output = X @ self.params['W1'].T + self.params['b1']
l1_relu = relu(l1_output)
l2_output = l1_relu @ self.params['W2'].T + self.params['b2']

y_pred = l2_output

```

ii.

2. FC

a. Affine Forward:

```

# ===== #
# YOUR CODE HERE:
# Calculate the output of the forward pass. Notice the dimensions
# of w are D x M, which is the transpose of what we did in earlier
# assignments.
# ===== #
N, D = x.shape[0], w.shape[0]
x_reshape = x.reshape(N, D)
out = x_reshape @ w + b

# ===== #
# END YOUR CODE HERE
# ===== #

cache = (x, w, b)
return out, cache

```

i.

b. Affine Backward:

```

N, D = x.shape[0], w.shape[0]
x_reshape = x.reshape(N, D)
dx = (dout @ w.T).reshape(x.shape)
dw = x_reshape.T @ dout
db = np.mean(np.ones((b.shape[0], dout.shape[0])) @ dout, axis=0)

# ===== #
# END YOUR CODE HERE
# ===== #

```

i. `return dx, dw, db`

c. RELU Forward:

```

"""
# ===== #
# YOUR CODE HERE:
#   Implement the ReLU forward pass.
# ===== #

def relu(a):
    result = a * (a > 0)
    return result

out = relu(x)
# ===== #
# END YOUR CODE HERE
# ===== #

cache = x
return out, cache

```

i.

d. RELU Backward:

```

"""
x = cache

# ===== #
# YOUR CODE HERE:
#   Implement the ReLU backward pass
# ===== #

# ReLU directs linearly to those > 0

dx = dout * (x > 0)

# ===== #
# END YOUR CODE HERE
# ===== #

```

i. `return dx`